

110612117 張仲瑜 HW1 report

I. Introduction

1. Goal

In this task, we aim to achieve as higher accuracy as possible in a 100-class image classification competition, with Resnet based CNN models.

2. Environment

- ✧ Laptop: Windows / RTX4060 * 1
- ✧ Kaggle account * 3: Linux / Tesla T4 * 2

3. Core Idea

I conducted a series of experiments to test whether designs such as [Cardinality, Squeeze-and-Excitation and various data augmentation](#) can boost the performance compared to original Resnet.

II. Method

1. Data augmentation set [1]

- ✧ Gaussian Blur + RandomResizeCrop

Since the size of the training images is not identical, we have to perform the down sampling on images to fit most of the pretrained models. To [avoid aliasing and mitigating noises](#), Gaussian Blur is applied before RandomResizeCrop.

- ✧ (Gaussian Blur) + Sharpening

After observing the dataset, I found that it's close to a multi-class (flower, grass, bird) fine-grained classification task. In my experience, [increasing the sharpness](#) of the data will help the model to identify [tiny differences](#) between them.

✧ Color Jittering + Greyscale

Since lots of grass and flowers are similar in colors, Color Jittering and Greyscale are applied to help the model [learning features except for the color](#).

✧ Mixup training [2]

Severe overfitting happens when finetuning a deep network, the [label-mixing strategy](#) increases the difficulty of training data, preventing the training loss from reducing to tiny numbers within few epochs

2. Model Architecture

✧ Model name: SEresNeXt101d_32x8d

Output size	SE-ResNeXt-101 (32×8d)		
144 × 144	conv, 7 × 7, 64, stride 2		
72 * 72	max pool, 3 × 3, stride 2		
72 × 72	conv, 1 × 1, 256 conv, 3 × 3, 256 conv, 1 × 1, 256 SE module	C=32	x3
36 x 36	conv, 1 × 1, 512 conv, 3 × 3, 512 conv, 1 × 1, 512 SE module	C=32	x4
18x18	conv, 1 × 1, 1024 conv, 3 × 3, 1024 conv, 1 × 1, 1024 SE module	C=32	x23
9x9	conv, 1 × 1, 2048 conv, 3 × 3, 2048 conv, 1 × 1, 2048 SE module	C=32	x3
1x1	dropout (p = 0.5)		
1x1	global average pooling, 100-d fc, softmax		

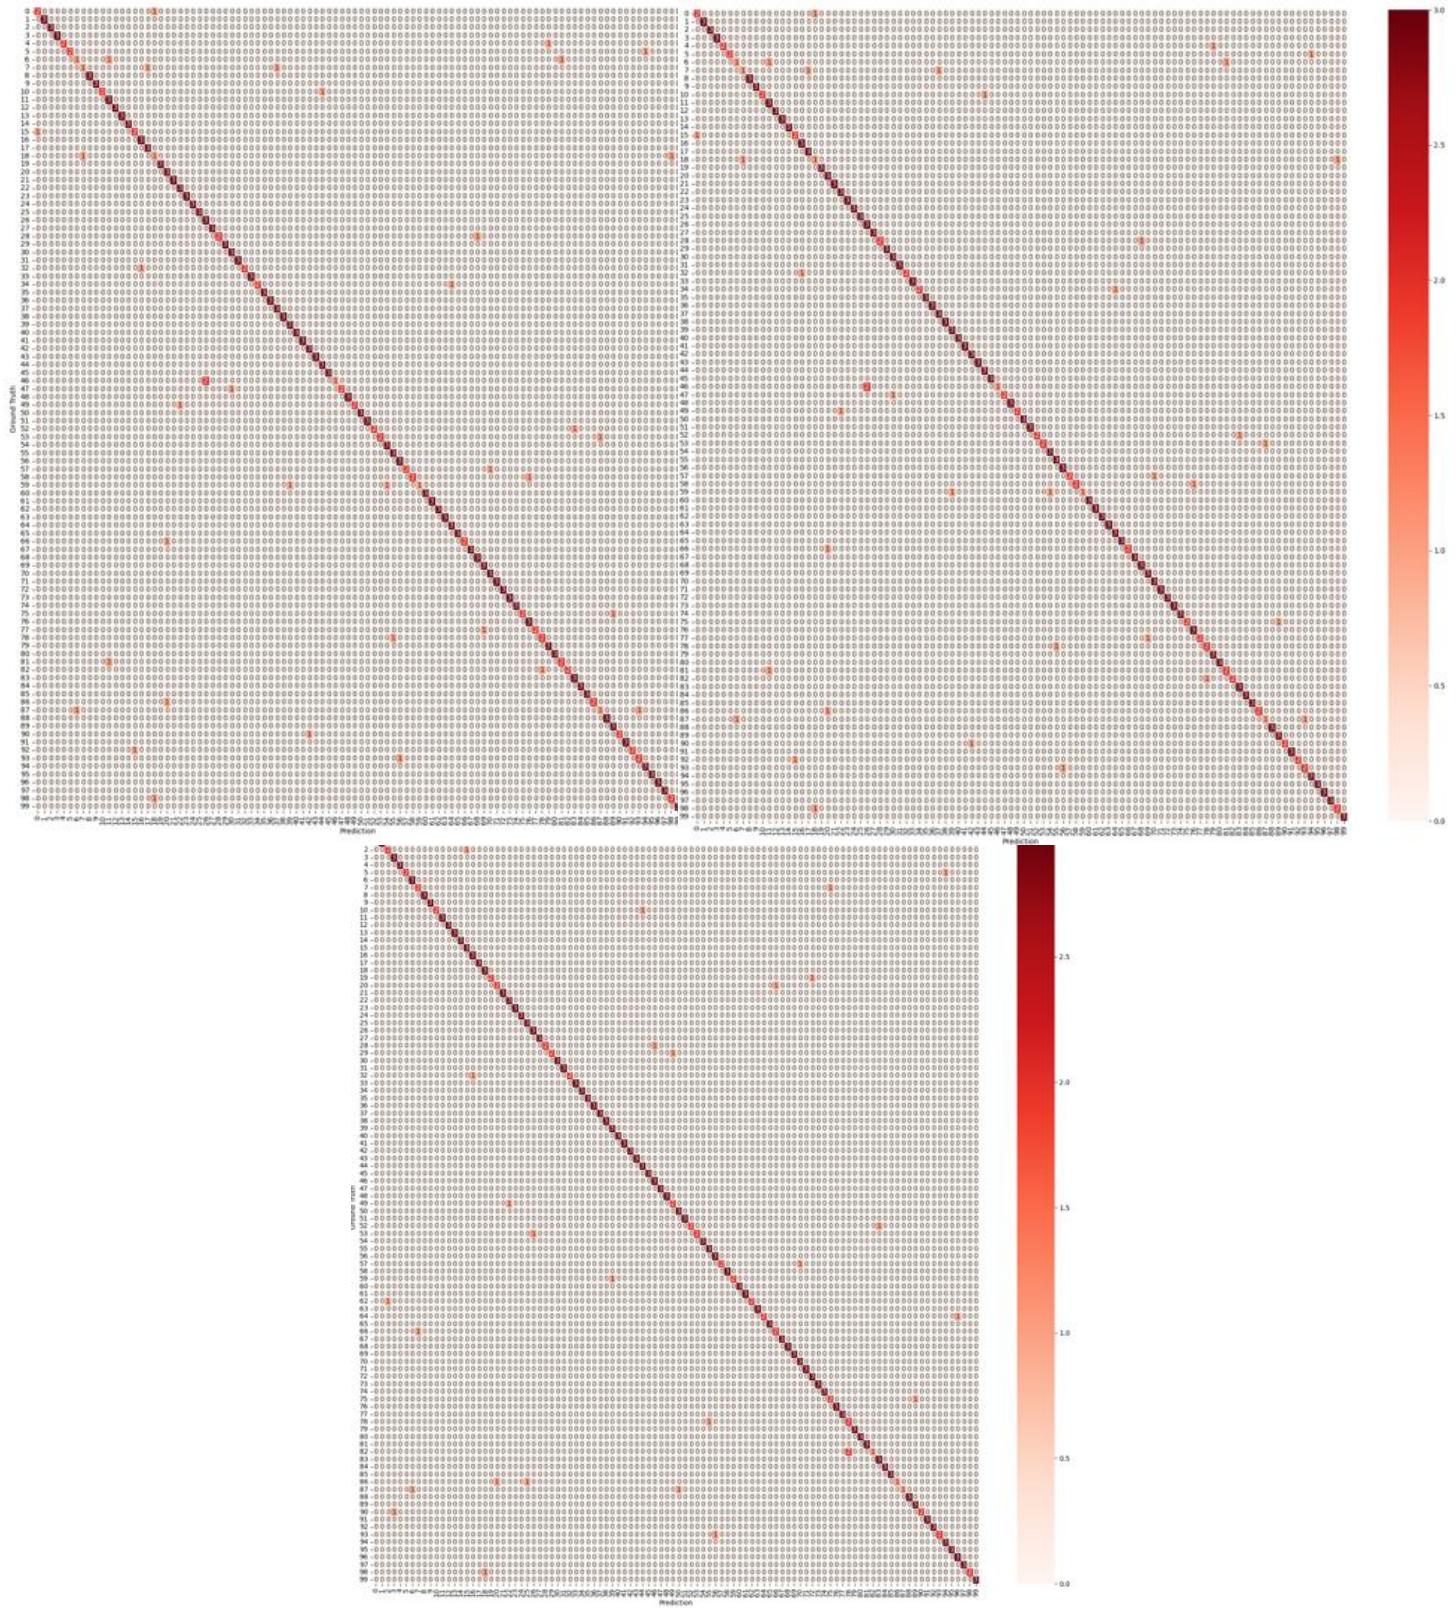
Total parameters: [93.6 M](#)

3. Hyperparameters (best one)
 - ✧ Pretrained on imagenet 1k
 - ✧ Training size: 224 * 224
 - ✧ Inferencing size: 320 * 320
 - ✧ Learning rate: Linear 10 warm-up [3] epochs to 5e-5 and * 0.5 decay factor whenever both validation accuracy and loss didn't decrease in 2 consecutive epochs.
Start from epoch 27, I reset the learning rate = 1.248e-6 with the same decay factor since I observe there exists a extreme point.
 - ✧ Loss: cross entropy with counter-data unbalance weight
 - ✧ Optimizer: AdamW
 - ✧ Batch size: 32, [4]
 - ✧ Epoch: Max is 100 but with early-stopping patience = 7 epochs, best epoch = 29
 - ✧ Weight decay = 1e-3
 - ✧ Mix-up probability = 0.5, alpha = 1.3

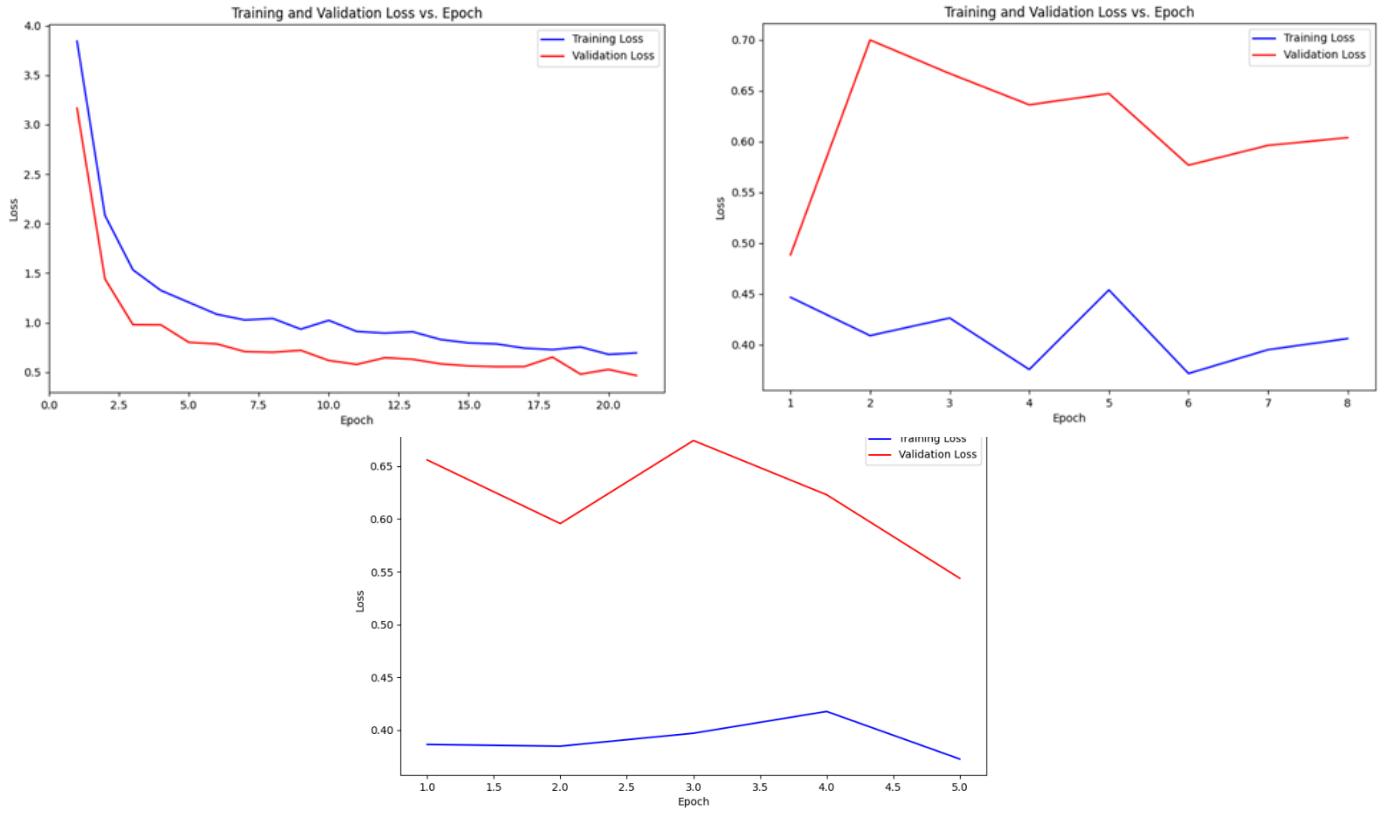
III. Result:

1. Best performance
 - ✧ Accuracy: [0.9633 on private testing](#)
Validation Loss: 0.6744, Validation Accuracy: 0.8933,
Validation Precision: 0.9155, [Validation Recall: 0.8933](#),
Validation F1: 0.8891

- ❖ Confusion Matrix (It breaks into three parts, I use strategy of continuing training with the same scalar, weight and optimizer state... in the first two one, and manually adjust the learning rate in the final state.):



✧ Learning Curve (loss):



IV. Additional experiment:

1. Hypothesis: Group Convolution is effective. The original paper claims that it can boost the performance of the model, so I want to verify it.

✧ Concept of Group Convolution [5][6]:

The idea is similar to the inception module, using **splitting**, **transforming and aggregating** to reduce the error with the same number of parameters.

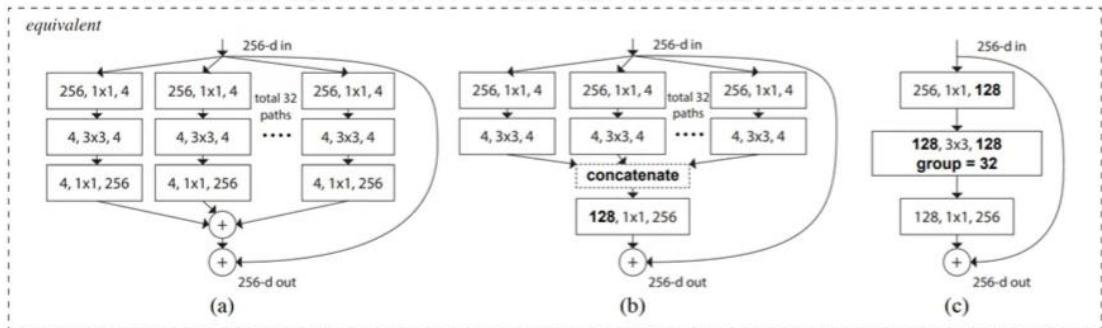


Fig from original paper

✧ Control variables:

Training size = Inferencing size = 224 * 224

Batch size = 20, dropout rate = 0.5, data augmentation set is used, learning rate = 1e-3, lr decay factor = 1e-3, no warm-up, no mix-up, weight decay = 1e-3

✧ Independent variable:

Cardinality in every bottleneck block

✧ Result

Model Architecture	Accuracy
resnet50 with 1x64d	0.86
resneXt50 with 32x4d	0.88

✧ Implication

As the author claims, the group convolution can boost model performance

2. Hypothesis: SE module is effective. The original paper claims that it can boost the performance of the model, so I want to verify it.

✧ Concept of SE [7][8]:

Adding a squeeze and excitation block (pooling -> fc -> non-linear activation function) at the end of every bottleneck block on top of the original resnet model, aiming to achieve channel-weight recalibration, aiming to let the model learn the relationship between channels and adjust their weight.

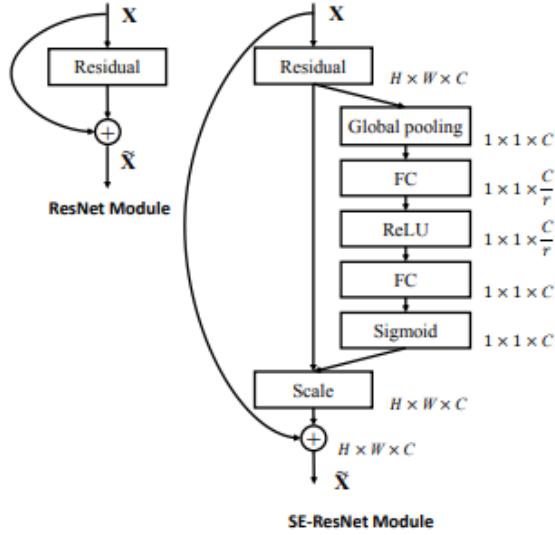


Fig from original paper

✧ Control variables:

Training size = Inferencing size = $224 * 224$

Batch size = 20, dropout rate = 0.5, data augmentation set is used, learning rate = $1e-3$, lr decay factor = $1e-3$, no warm-up, no mix-up, weight decay = $1e-3$

✧ Independent variable:

SE block in every bottleneck block

✧ Result

Model Architecture	Accuracy
resneXt50_324D w/o SE	0.88
resneXt50_324D with SE	0.91

✧ Implication

As the author of the claims, adding SE blocks is significantly effective.

3. Hypothesis: Using reciprocal of class-number as its weight in the Cross Entropy loss can improve the performance

✧ Concept: Through my observation, I found that the dataset is unbalanced. Comparing two models with similar loss, the one with a higher recall rate performs better on the coda bench, so I set classes weight in CE as the reciprocal of their number of data, aiming to solve this problem.

✧ Result

Loss formula	Accuracy
average weighted	0.9570
reciprocal version	0.9633

✧ Implication:
It works, dealing with those classes is important!

V. Github

https://github.com/sharkccy/NYCU_CV_2025_Spring

VI. Thoughts

I learned more about how to run a huge model within limited resources. At the very beginning, my poor code can only run a resnet50 model and taking 10 minute per epoch. I thought running a SEresneXt101 would be impossible, it turned out that properly garbage collection, CUDA management, properly num_workers and closing plots are key during training. It's an unforgettable experience to witness "threshing" happened in 2025. After those adjustments, I can run SEresneXt101 in 15 minutes per epoch.

model	epoch	Drop out	loss	optimizer	Data augmentation set	coda bench	learning rate/decay factor	lr_patience	training size	testing size	batch size	warm_up	mix up
resnet18	183/200	N	CE	AdamW	Y	N/A	1e-3	3	224	224	64	N	N
resnet50	20	0.5	CE	AdamW	N	0.86	1e-3~1e-6	3	224	224	64	N	N
resnet101	20	0.5	CE	AdamW	N	0.85	1e-3~1e-5	3	224	224	64	N	N
resneXt50_324D	20	0.5	CE	AdamW	N	0.88	1e-3~1e-6	3	224	224	64	N	N
SEresneXt50_324D	20	0.5	CE	AdamW	sharpening	0.86	1e-3~1e-6	3	224	224	64	N	N
SEresneXt50_324D	20	0.5	CE	AdamW	N	0.88	1e-3~1e-5	3	224	224	64	N	N
SEresneXt50_324D	20	0.5	CE	AdamW	N	0.9	1e-3~1e-6	2	224	224	64	N	N
SEresneXt50_324D	20	0.5	CE	AdamW	N	0.91	1e-3~1e-6	2	224	288	64	N	N
SEresneXt50_324D	20	0.2	CE	AdamW	Y	0.9	1e-3~1e-6	2	320	320	32	N	N
seresnet152.ra2_in1k	28 / 50	0.5	CE	AdamW	Y(new)	0.93	1e-5 / 0.8	2	256	320	32	Y	N
seresnet152.ra2_in1k	20 / 50	0.5	CE	AdamW	N	0.9	1e-3~1e-6	2	256	320	32	N	N
seresnet152.ra2_in1k	41 / 50	0.5	CE	AdamW	Y(new)	0.91	1e-3 / 0.8 / 39 後 0.5	2	256	320	32	Y	0.5 / 1.3
seresnet152.ra2_in1k	/ 50	0.5	CE	AdamW	Y(new)		1e-3 / 0.8	2	256	320	32	N	0.5 / 1.3
seresnet152.ra2_in1k	20 / 50												
resnet152.ra2_in1k	/ 100												
seresnextaa101d_32x8d	28 / 100	0.5	CE	AdamW	Y(New)	0.96	5e-5 / 0.5	1	224	320	32	N	1 / 1.3
seresnextaa101d_32x8d	100	0.5	CE(W)	AdamW	Y(New)		5e-5 / 0.5	1	224	320	32	3	1 / 1.3
seresnextaa101d_32x8d	21 / 35	0.5	CE(W)	AdamW	Y(New)	0.96	5e-5 / 0.5	1	224	320	32	5	0.5 / 1.3
seresnextaa101d_32x8d	30 / 100	0.5	CE(W)	AdamW	Y(New)	0.96	5e-5 / 0.5, 1.2	1	224	320	32	10	0.5 / 1.3
seresnextaa101d_32x8d	30	0.5	CE	AdamW	Y(New)	0.96	5e-5 / 0.5	1	224	320	32	3	N
seresnextaa101d_32x8d	23 / 25	0.5	CE	AdamW	Y(New)	0.96	5e-5 / 0.5	1	224	320	32	N	N
seresnextaa101d_32x8d	25	0.5	CE	AdamW	Y(New)	0.96	5e-5 / 0.5	1	224	288	32	N	N
resnetrs50_tf_in1k	13	0.2	CE	AdamW	Y	0.88			224	224	32	N	

Fig. Some of my experiments

VII. Reference:

1. <https://pytorch.org/vision/main/transforms.html>
2. <https://blog.csdn.net/xiaoxifei/article/details/90416997>
3. <https://reurl.cc/Z4jjkM>
4. <https://reurl.cc/QY66AM>
5. <https://liaowc.github.io/blog/resnext-structure/>
6. <https://reurl.cc/W0XX9L>
7. <https://meetonfriday.com/posts/79fdff34/>
8. <https://liaowc.github.io/blog/SENet-structure/>
9. <https://github.com/huggingface/pytorch-image-models/blob/main/results/results-imagenet.csv>