

110612117 張仲瑜 HW2 report

I. Introduction

1. Goal

In this task, we aim to achieve as high mAP and accuracy as possible in a digit detection competition with Faster-RCNN based model. There are two tasks in the competition, including digit bounding box prediction and digits recognition in images prediction.

2. Environment

- ✧ Laptop: Windows / RTX4060 * 1
- ✧ Kaggle account * 3: Linux / Tesla T4 * 2

3. Core Idea

I built a fast RCNN pipeline [1,2] with custom components, including backbone, anchor generator [3] and ROI_pooler, which allows me to select a pretrained ResNet101d as the backbone model while designing the FPN (Anchor generater as I want to conduct experiments.

II. Method

1. Data augmentation set

✧ Resize

After observing the data that are confusing to the model, which are either small or low resolution, I perform the upsampling with bicubic interpolation to the data and adjust the label simultaneously to increase the mAP [4] in task one.

✧ Contrast Enhancement, Sharpening

Some of the data are difficult since lights are dim in those images, I perform the upsampling with bicubic interpolation

to increase the performance of the pipeline.

2. Model Architecture

- ✧ Backbone Model: ResNet101 (Pretrained on imageNet1k)

Output size	ResNet-101
144×144	conv, 7×7 , 64, stride 2
72×72	max pool, 3×3 , stride 2
72×72	conv, 1×1 , 256
	conv, 3×3 , 256
	conv, 1×1 , 256
36×36	conv, 1×1 , 512
	conv, 3×3 , 512
	conv, 1×1 , 512
18×18	conv, 1×1 , 1024
	conv, 3×3 , 1024
	conv, 1×1 , 1024
9×9	conv, 1×1 , 2048
	conv, 3×3 , 2048
	conv, 1×1 , 2048
1×1	global average pooling, 2048-d fc, softmax

- ✧ FPN: I use all four 4 layers of the model

```
def forward(self, x):
    x = self.stem(x)
    c2 = self.layer1(x)
    c3 = self.layer2(c2)
    c4 = self.layer3(c3)
    c5 = self.layer4(c4)
    return {
        '0': c2,
        '1': c3,
        '2': c4,
        '3': c5
    }
```

- ✧ Anchor size map and aspect ratio:
Ascending pixels size as layers go deeper, with the common digit aspect ratio.

```
anchor_generator = torchvision.models.detection.rpn.AnchorGenerator(
    sizes = ((4, 8, 16), (8, 16, 32), (16, 32, 64), (32, 64, 128), (64, 128, 256)),
    # sizes = ((4, 12, 20), (28, 36, 44), (52, 60, 68), (68, 72, 80), (100, 128, 256)),
    aspect_ratios = ((0.5, 1.0, 2.0),) * 5,
)
```

- ✧ ROI pooler:
4 feature maps with default settings.

```
roi_pooler = torchvision.ops.MultiscaleRoIAlign(
    featmap_names=['0', '1', '2', '3'],
    output_size=7,
    sampling_ratio=2
)
```

3. Hyperparameters (best one)
 - ✧ Pretrained on imageNet1k
 - ✧ Training/ testing box NMS threshold [5]: 0.5
 - ✧ Training/ testing box score threshold: 0.65
 - ✧ Anchor map: As image above
 - ✧ Aspect ratios: As image above
 - ✧ Sharpening factor = 1.5
 - ✧ Learning rate: 5e-5 and * 0.5 decay factor whenever both validation accuracy and loss didn't decrease in 2 consecutive epochs.
 - ✧ Loss: Cross Entropy (classification) + Smooth L1(Regression)
 - ✧ Optimizer: AdamW
 - ✧ Batch size: 4
 - ✧ Epoch: Max is 50 but with early-stopping patience = 7 epochs, stopping at epoch 28.
 - ✧ Weight decay = 1e-3

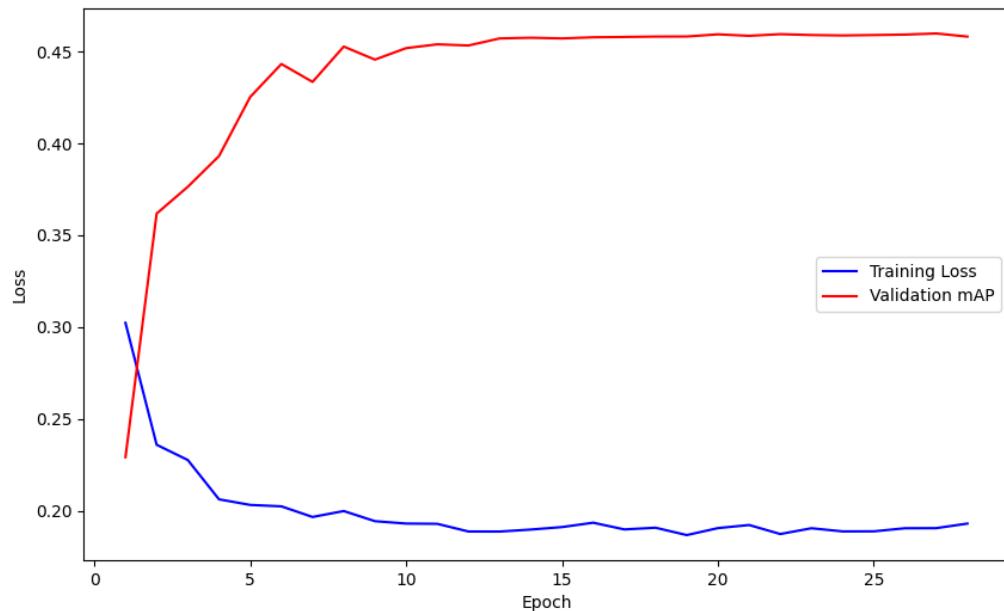
III. Result:

1. Best performance

- ✧ Task 1 mAP: **0.3767** on private testing
- ✧ Task 2 Accuracy: **0.8477** on private testing
- ✧ Other COCO indicators

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.456
Average Precision (AP) @[ IoU=0.50    | area= all | maxDets=100 ] = 0.907
Average Precision (AP) @[ IoU=0.75    | area= all | maxDets=100 ] = 0.384
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.446
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.497
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.597
Average Recall   (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.504
Average Recall   (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.537
Average Recall   (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.537
Average Recall   (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.527
Average Recall   (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.572
Average Recall   (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.604
```

- ✧ Learning curve (epoch vs Training loss & Val mAP)



IV. Additional experiment [6]:

1. Hypothesis: Adjust the custom anchor size stride in each feature map to 4, that is ((4, 12, 20), (28, 36, 44), (52, 60, 68), (68, 72, 80), (100, 128, 256)), instead of double one ((4, 8, 16)...) will increase the performance on small targets.
 - ✧ I think this may work because most of our digit sizes lie in a specific dense range.
 - ✧ Result

AP@IoU=0.50:0.95	
Anchor stride	area = small
Stride = +4	0.446
Stride = x2	0.443

- ✧ Implication

There's only a slight difference between them; I think that's because in latter layers, the receptive fields are too large that the influence of particular small area is little, that is, the boost only comes from the stride change in first two layers.

2. Hypothesis: The higher the scale factor in supe resolution is, the better the performance is.

✧ I think this may work after observing images that are difficult to recognize, and most of them are small.

✧ Result

AP@IoU=0.50:0.95	
Scale Factor	area = small
4	0.446
2	0.438

✧ Implication

It seems that under the same settings, larger scale factors did help the model learn slightly better in small areas; however, it also increased lots of computational cost.

Personally, I think it's not worth doing it.

3. Hypothesis: The higher the box NMS threshold is, the more digits we will get in the prediction, which means that we can observe the wrong data type (missing/repeat digits) to adjust the threshold.

- ✧ I think this may work since NMS will remove boxes that have high IoU with the one which has the highest confident score. Sometimes numbers like 117 or 112 will be recognized as 17, 1117, 12 and 1112 depending on this threshold
- ✧ Result

Box_NMS_threshold	result
0.7	7117/1112
0.5	117/112
0.3	117/112

- ✧ Implication:
It did work and I successfully boost the performance through adjusting the threshold to train a better model.

V. Github

https://github.com/sharkccy/NYCU_CV_2025_Spring/tree/main/hw2

VI. Thoughts

It's my first time running object detection models locally, which makes me learn more about how to build a Faster RCNN pipeline given CNN model. The process of observing and barnstorming is interesting as usual, and I think data processing is still the key in this kind of task. Having great datasets is far more effective than other training strategies.

1	settings	Average Precision (AP) @[IoU=0.50:0.95 area= all maxDets=100] = 0.458	Average Precision (AP) @[IoU=0.50:0.95 area= all maxDets=100] = 0.455	Average Precision (AP) @[IoU=0.50:0.95 area= all maxDets=100] = 0.455
2	0703normal_anchor	Average Precision (AP) @[IoU=0.50 area= all maxDets=100] = 0.001	Average Precision (AP) @[IoU=0.50 area= all maxDets=100] = 0.904	Average Precision (AP) @[IoU=0.50 area= all maxDets=100] = 0.898
3	0703normal_noaug	Average Precision (AP) @[IoU=0.75 area= all maxDets=100] = 0.402	Average Precision (AP) @[IoU=0.75 area= all maxDets=100] = 0.384	Average Precision (AP) @[IoU=0.75 area= all maxDets=100] = 0.389
4	0703_4_anchor	Average Precision (AP) @[IoU=0.50:0.95 area= small maxDets=100] = 0.448	Average Precision (AP) @[IoU=0.50:0.95 area= small maxDets=100] = 0.445	Average Precision (AP) @[IoU=0.50:0.95 area= small maxDets=100] = 0.443
5	0703_8_anchor	Average Precision (AP) @[IoU=0.50:0.95 area= medium maxDets=100] = 0.499	Average Precision (AP) @[IoU=0.50:0.95 area= medium maxDets=100] = 0.497	Average Precision (AP) @[IoU=0.50:0.95 area= medium maxDets=100] = 0.499
6	05_065_4_anchor	Average Precision (AP) @[IoU=0.50:0.95 area= large maxDets=100] = 0.590	Average Precision (AP) @[IoU=0.50:0.95 area= large maxDets=100] = 0.565	Average Precision (AP) @[IoU=0.50:0.95 area= large maxDets=100] = 0.564
7	05_065_8_anchor	Average Recall (AR) @[IoU=0.50:0.95 area= all maxDets= 10] = 0.541	Average Recall (AR) @[IoU=0.50:0.95 area= all maxDets= 10] = 0.536	Average Recall (AR) @[IoU=0.50:0.95 area= all maxDets= 10] = 0.564
8		Average Recall (AR) @[IoU=0.50:0.95 area= small maxDets=100] = 0.541	Average Recall (AR) @[IoU=0.50:0.95 area= small maxDets=100] = 0.526	Average Recall (AR) @[IoU=0.50:0.95 area= small maxDets=100] = 0.557
9	05_065_4x065	Average Recall (AR) @[IoU=0.50:0.95 area= medium maxDets=100] = 0.531	Average Recall (AR) @[IoU=0.50:0.95 area= medium maxDets=100] = 0.571	Average Recall (AR) @[IoU=0.50:0.95 area= medium maxDets=100] = 0.589
10	05_065_3x065	Average Recall (AR) @[IoU=0.50:0.95 area= large maxDets=100] = 0.601	Average Recall (AR) @[IoU=0.50:0.95 area= large maxDets=100] = 0.582	Average Recall (AR) @[IoU=0.50:0.95 area= large maxDets=100] = 0.605
11	05_065_2x065			
12	05_065_1x065	' Average Precision (AP) @[IoU=0.50:0.95 area= all maxDets=100] = 0.459	7 Average Precision (AP) @[IoU=0.50:0.95 area= all maxDets=100] = 0.458	Average Precision (AP) @[IoU=0.50:0.95 area= all maxDets=100] = 0.455
13	05_065_1x055	' Average Precision (AP) @[IoU=0.50 area= all maxDets=100] = 0.900	8 Average Precision (AP) @[IoU=0.50 area= all maxDets=100] = 0.896	Average Precision (AP) @[IoU=0.50 area= all maxDets=100] = 0.898
14	05_065_1x050	' Average Precision (AP) @[IoU=0.75 area= all maxDets=100] = 0.401	9 Average Precision (AP) @[IoU=0.75 area= all maxDets=100] = 0.400	Average Precision (AP) @[IoU=0.75 area= all maxDets=100] = 0.389
15		' Average Precision (AP) @[IoU=0.50:0.95 area= small maxDets=100] = 0.449	10 Average Precision (AP) @[IoU=0.50:0.95 area= small maxDets=100] = 0.446	Average Precision (AP) @[IoU=0.50:0.95 area= small maxDets=100] = 0.443
16		' Average Precision (AP) @[IoU=0.50:0.95 area= medium maxDets=100] = 0.502	11 Average Precision (AP) @[IoU=0.50:0.95 area= medium maxDets=100] = 0.504	Average Precision (AP) @[IoU=0.50:0.95 area= medium maxDets=100] = 0.499
17		' Average Precision (AP) @[IoU=0.50:0.95 area= large maxDets=100] = 0.598	12 Average Precision (AP) @[IoU=0.50:0.95 area= large maxDets=100] = 0.594	Average Precision (AP) @[IoU=0.50:0.95 area= large maxDets=100] = 0.598
18		' Average Recall (AR) @[IoU=0.50:0.95 area= all maxDets= 10] = 0.505	13 Average Recall (AR) @[IoU=0.50:0.95 area= all maxDets= 10] = 0.503	Average Recall (AR) @[IoU=0.50:0.95 area= all maxDets= 10] = 0.500
19		' Average Recall (AR) @[IoU=0.50:0.95 area= all maxDets=100] = 0.567	14 Average Recall (AR) @[IoU=0.50:0.95 area= all maxDets=100] = 0.569	Average Recall (AR) @[IoU=0.50:0.95 area= all maxDets=100] = 0.564
20		' Average Recall (AR) @[IoU=0.50:0.95 area= small maxDets=100] = 0.559	15 Average Recall (AR) @[IoU=0.50:0.95 area= small maxDets=100] = 0.560	Average Recall (AR) @[IoU=0.50:0.95 area= small maxDets=100] = 0.557
21		' Average Recall (AR) @[IoU=0.50:0.95 area= medium maxDets=100] = 0.598	16 Average Recall (AR) @[IoU=0.50:0.95 area= medium maxDets=100] = 0.571	Average Recall (AR) @[IoU=0.50:0.95 area= medium maxDets=100] = 0.589
22		' Average Recall (AR) @[IoU=0.50:0.95 area= large maxDets=100] = 0.622	17 Average Recall (AR) @[IoU=0.50:0.95 area= large maxDets=100] = 0.637	Average Recall (AR) @[IoU=0.50:0.95 area= large maxDets=100] = 0.605

Fig. Some of my experiments

VII. Reference:

1. <https://ivan-eng-murmur.medium.com/object-detection-s3-faster-rcnn-%E7%B0%A1%E4%BB%8B-5f37b13ccdd2>
2. <https://didi3310781.medium.com/faster-r-cnn-%E5%AD%B8%E7%BF%92%E7%AD%86%E8%A8%98-4917d59edb55>
3. <https://github.com/pytorch/vision/issues/3246>
4. <https://reurl.cc/Dq48ee>
5. <https://reurl.cc/0K9yn9>
6. https://pytorch.org/vision/main/_modules/torchvision/models/detection/faster_rcnn.html#fasterrcnn_resnet50_fpn_v2